

```

import pandas as pd
import numpy as np
import missingno as msno
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots
import warnings
warnings.filterwarnings('ignore')

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from xgboost import XGBClassifier

from sklearn import metrics
from sklearn.metrics import roc_curve
from sklearn.metrics import recall_score, confusion_matrix, precision_score, f1_score, accuracy_score, classification_report

df=pd.read_csv('/content/WA_Fn-UseC_-Telco-Customer-Churn.csv')

```

```
df.head()
```



	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultiLine
0	7590-VHVEG	Female	0	Yes	No	1	No	No
1	5575-GNVDE	Male	0	No	No	34	Yes	Yes
2	3668-QPYBK	Male	0	No	No	2	Yes	Yes
3	7795-CFOCW	Male	0	No	No	45	No	No
4	9237-HQITU	Female	0	No	No	2	Yes	Yes

5 rows × 21 columns

```
df.shape
```



(7043, 21)

```
df.info()
```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   customerID            7043 non-null  object  
1   gender                 7043 non-null  object  
2   SeniorCitizen          7043 non-null  int64   
3   Partner                7043 non-null  object  
4   Dependents             7043 non-null  object  
5   tenure                 7043 non-null  int64   
6   PhoneService           7043 non-null  object  
7   MultipleLines          7043 non-null  object  
8   InternetService        7043 non-null  object  
9   OnlineSecurity         7043 non-null  object  
10  OnlineBackup           7043 non-null  object  
11  DeviceProtection       7043 non-null  object  
12  TechSupport            7043 non-null  object  
13  StreamingTV            7043 non-null  object  

```

```

14 StreamingMovies    7043 non-null    object
15 Contract           7043 non-null    object
16 PaperlessBilling   7043 non-null    object
17 PaymentMethod      7043 non-null    object
18 MonthlyCharges     7043 non-null    float64
19 TotalCharges       7043 non-null    object
20 Churn              7043 non-null    object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

```
df.columns.values
```

```

array(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
      'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
      'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges',
      'TotalCharges', 'Churn'], dtype=object)

```

```
df.dtypes
```

```

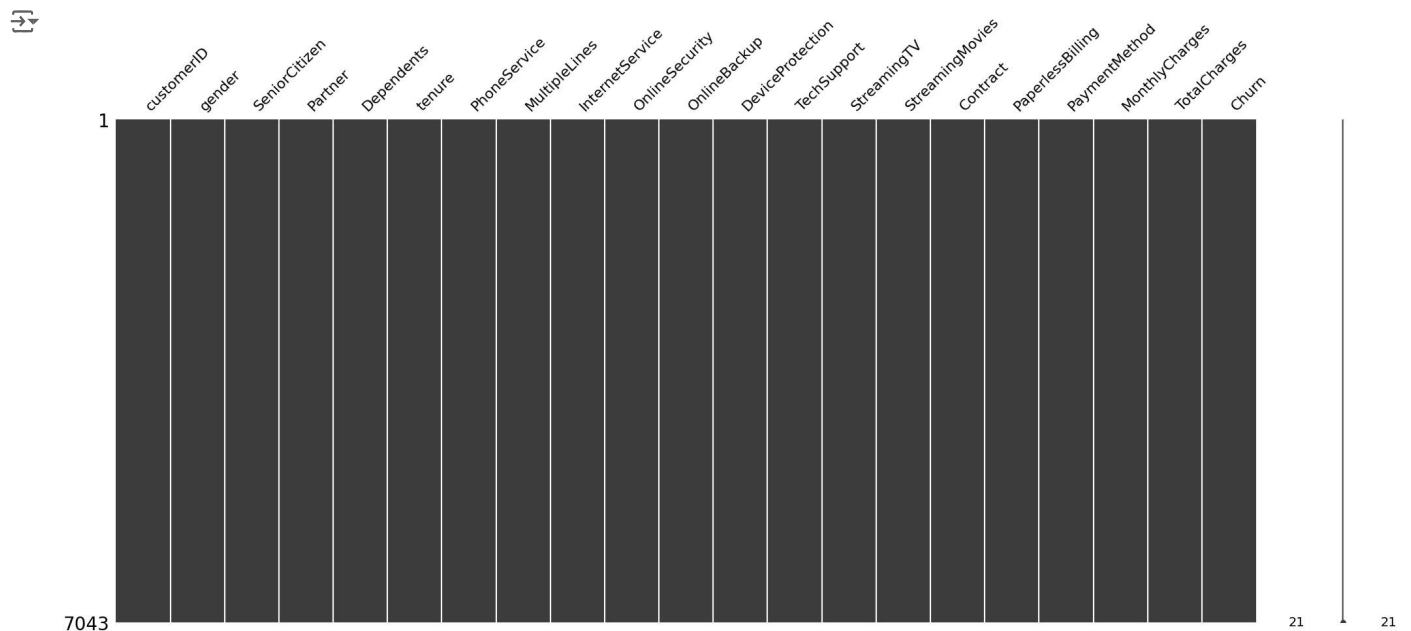
customerID      object
gender          object
SeniorCitizen   int64
Partner         object
Dependents      object
tenure          int64
PhoneService    object
MultipleLines   object
InternetService object
OnlineSecurity  object
OnlineBackup    object
DeviceProtection object
TechSupport     object
StreamingTV     object
StreamingMovies object
Contract        object
PaperlessBilling object
PaymentMethod   object
MonthlyCharges  float64
TotalCharges    object
Churn           object
dtype: object

```

```

# Visualize missing values as a matrix
msno.matrix(df);

```



```

df = df.drop(['customerID'], axis = 1)
df.head()

```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	D
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	
1	Male	0	No	No	34	Yes	No	DSL	Yes	No	
2	Male	0	No	No	2	Yes	No	DSL	Yes	Yes	
3	Male	0	No	No	45	No	No phone service	DSL	Yes	No	
4	Female	0	No	No	2	Yes	No	Fiber optic	No	No	

Next steps: [Generate code with df](#) [View recommended plots](#)

```
df['TotalCharges'] = pd.to_numeric(df.TotalCharges, errors='coerce')
df.isnull().sum()
```

```
gender      0
SeniorCitizen  0
Partner      0
Dependents   0
tenure       0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport   0
StreamingTV    0
StreamingMovies  0
Contract       0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  11
Churn          0
dtype: int64
```

```
df[np.isnan(df['TotalCharges'])]
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	D
488	Female	0	Yes	Yes	0	No	No phone service	DSL	Yes	No	
753	Male	0	No	Yes	0	Yes	No	No	No internet service	No internet service	
936	Female	0	Yes	Yes	0	Yes	No	DSL	Yes	Yes	
1082	Male	0	Yes	Yes	0	Yes	Yes	No	No internet service	No internet service	
1340	Female	0	Yes	Yes	0	No	No phone service	DSL	Yes	Yes	
3331	Male	0	Yes	Yes	0	Yes	No	No	No internet service	No internet service	
3826	Male	0	Yes	Yes	0	Yes	Yes	No	No internet service	No internet service	
4380	Female	0	Yes	Yes	0	Yes	No	No	No internet service	No internet service	
5218	Male	0	Yes	Yes	0	Yes	No	No	No internet service	No internet service	
6670	Female	0	Yes	Yes	0	Yes	Yes	DSL	No	Yes	
6754	Male	0	No	Yes	0	Yes	Yes	DSL	Yes	Yes	

```
df[df['tenure'] == 0].index
```

```
Index([488, 753, 936, 1082, 1340, 3331, 3826, 4380, 5218, 6670, 6754], dtype='int64')
```

```
df.drop(labels=df[df['tenure'] == 0].index, axis=0, inplace=True)
df[df['tenure'] == 0].index
```

➦ Index([], dtype='int64')

```
df.fillna(df["TotalCharges"].mean())
```

➦

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes
1	Male	0	No	No	34	Yes	No	DSL	Yes	No
2	Male	0	No	No	2	Yes	No	DSL	Yes	Yes
3	Male	0	No	No	45	No	No phone service	DSL	Yes	No
4	Female	0	No	No	2	Yes	No	Fiber optic	No	No
...
7038	Male	0	Yes	Yes	24	Yes	Yes	DSL	Yes	No
7039	Female	0	Yes	Yes	72	Yes	Yes	Fiber optic	No	Yes
7040	Female	0	Yes	Yes	11	No	No phone service	DSL	Yes	No
7041	Male	1	Yes	No	4	Yes	Yes	Fiber optic	No	No
7042	Male	0	No	No	66	Yes	No	Fiber optic	Yes	No

7032 rows × 20 columns

```
df.isnull().sum()
```

➦

gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	0
Churn	0

dtype: int64

```
df["SeniorCitizen"] = df["SeniorCitizen"].map({0: "No", 1: "Yes"})
df.head()
```

➦

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	D
0	Female	No	Yes	No	1	No	No phone service	DSL	No	Yes	
1	Male	No	No	No	34	Yes	No	DSL	Yes	No	
2	Male	No	No	No	2	Yes	No	DSL	Yes	Yes	
3	Male	No	No	No	45	No	No phone service	DSL	Yes	No	
4	Female	No	No	No	2	Yes	No	Fiber optic	No	No	

Next steps: [Generate code with df](#)

☒ [View recommended plots](#)

```
df["InternetService"].describe(include=['object', 'bool'])
```

```
count      7032
unique       3
top    Fiber optic
freq      3096
Name: InternetService, dtype: object
```

```
numerical_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']
df[numerical_cols].describe()
```


```
count      7032.000000    7032.000000    7032.000000
mean       32.421786      64.798208    2283.300441
std        24.545260      30.085974    2266.771362
min         1.000000      18.250000     18.800000
25%         9.000000      35.587500    401.450000
50%        29.000000      70.350000   1397.475000
75%        55.000000      89.862500   3794.737500
max        72.000000     118.750000   8684.800000
```

```
g_labels = ['Male', 'Female']
c_labels = ['No', 'Yes']
# Create subplots: use 'domain' type for Pie subplot
fig = make_subplots(rows=1, cols=2, specs=[[{'type':'domain'}], [{'type':'domain'}]])
fig.add_trace(go.Pie(labels=g_labels, values=df['gender'].value_counts(), name="Gender"),
              1, 1)
fig.add_trace(go.Pie(labels=c_labels, values=df['Churn'].value_counts(), name="Churn"),
              1, 2)
```

```
# Use `hole` to create a donut-like pie chart
fig.update_traces(hole=.4, hoverinfo="label+percent+name", textfont_size=16)
```

```
fig.update_layout(
    title_text="Gender and Churn Distributions",
    # Add annotations in the center of the donut pies.
    annotations=[dict(text='Gender', x=0.16, y=0.5, font_size=20, showarrow=False),
                  dict(text='Churn', x=0.84, y=0.5, font_size=20, showarrow=False)])
fig.show()
```

Gender and Churn Distributions



```
df["Churn"][df["Churn"]=="No"].groupby(by=df["gender"]).count()
```

```
gender
Female    2544
```

```
Male      2619
Name: Churn, dtype: int64
```

```
df["Churn"][df["Churn"]=="Yes"].groupby(by=df["gender"]).count()
```

```
gender
Female    939
Male      930
Name: Churn, dtype: int64
```


```
plt.figure(figsize=(6, 6))
labels = ["Churn: Yes", "Churn:No"]
values = [1869, 5163]
labels_gender = ["F", "M", "F", "M"]
sizes_gender = [939, 930, 2544, 2619]
colors = ['#ff6666', '#66b3ff']
colors_gender = ['#c2c2f0', '#ffb3e6', '#c2c2f0', '#ffb3e6']
explode = (0.3, 0.3)
explode_gender = (0.1, 0.1, 0.1, 0.1)
textprops = {"fontsize": 15}
#Plot
plt.pie(values, labels=labels, autopct='%1.1f%%', pctdistance=1.08, labeldistance=0.8, colors=colors, startangle=90, frame=True, explode=explode)
plt.pie(sizes_gender, labels=labels_gender, colors=colors_gender, startangle=90, explode=explode_gender, radius=7, textprops=textprops, counterclock=True)
#Draw circle
centre_circle = plt.Circle((0,0), 5, color='black', fc='white', linewidth=0)
fig = plt.gcf()
fig.gca().add_artist(centre_circle)

plt.title('Churn Distribution w.r.t Gender: Male(M), Female(F)', fontsize=15, y=1.1)

# show plot

plt.axis('equal')
plt.tight_layout()
plt.show()
```


 Churn Distribution w.r.t Gender: Male(M), Female(F)



```
fig = px.histogram(df, x="Churn", color="Contract", barmode="group", title="Customer contract distribution")
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```



Customer contract distribution




```
labels = df['PaymentMethod'].unique()
values = df['PaymentMethod'].value_counts()

fig = go.Figure(data=[go.Pie(labels=labels, values=values, hole=.3)])
fig.update_layout(title_text="Payment Method Distribution")
fig.show()
```




Payment Method Distribution



```
fig = px.histogram(df, x="Churn", color="PaymentMethod", title="Customer Payment Method distribution w.r.t. Churn")
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```



Customer Payment Method distribution w.r.t. Churn



```
df["InternetService"].unique()
```

```
array(['DSL', 'Fiber optic', 'No'], dtype=object)
```

```
df[df["gender"]=="Male"][["InternetService", "Churn"].value_counts()
```

```
InternetService  Churn
DSL              No    992
Fiber optic     No    910
No              No    717
Fiber optic     Yes   633
DSL             Yes   240
No              Yes    57
Name: count, dtype: int64
```

```
df[df["gender"]=="Female"][["InternetService", "Churn"].value_counts()
```

```
InternetService  Churn
DSL              No    965
Fiber optic     No    889
No              No    690
Fiber optic     Yes   664
DSL             Yes   219
No              Yes    56
Name: count, dtype: int64
```

```
fig = go.Figure()
```

```
fig.add_trace(go.Bar(
    x = [['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
        ["Female", "Male", "Female", "Male"]],
    y = [965, 992, 219, 240],
    name = 'DSL',
))
```

```
fig.add_trace(go.Bar(
    x = [['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
        ["Female", "Male", "Female", "Male"]],
    y = [889, 910, 664, 633],
    name = 'Fiber optic',
))
```


```
fig.add_trace(go.Bar(
    x = [['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
        ["Female", "Male", "Female", "Male"]],
    y = [690, 717, 56, 57],
    name = 'No Internet',
))
```

```
fig.update_layout(title_text="Churn Distribution w.r.t. Internet Service and Gender")
```

```
fig.show()
```





Churn Distribution w.r.t. Internet Service and Gender



```
color_map = {"Yes": "#FF97FF", "No": "#AB63FA"}
fig = px.histogram(df, x="Churn", color="Dependents", barmode="group", title="Dependents distribution", color_discrete_map=color_map)
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```



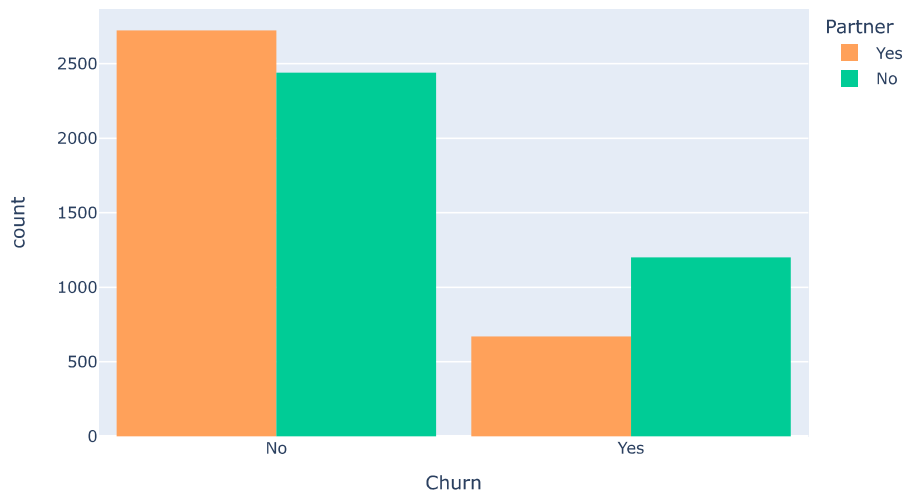
Dependents distribution



```
color_map = {"Yes": '#FFA15A', "No": '#00CC96'}
fig = px.histogram(df, x="Churn", color="Partner", barmode="group", title="Churn distribution w.r.t. Partners", color_discrete_map=color_map)
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```




Chrunch distribution w.r.t. Partners



```
color_map = {"Yes": '#00CC96', "No": '#B6E880'}  
fig = px.histogram(df, x="Churn", color="SeniorCitizen", title="Churn distribution w.r.t. Senior Citizen", color_discrete_map=co  
fig.update_layout(width=700, height=500, bargap=0.1)  
fig.show()
```




Churn distribution w.r.t. Senior Citizen



```
color_map = {"Yes": "#FF97FF", "No": "#AB63FA"}  
fig = px.histogram(df, x="Churn", color="OnlineSecurity", barmode="group", title="Churn w.r.t Online Security", color_discrete_m  
fig.update_layout(width=700, height=500, bargap=0.1)  
fig.show()
```




Churn w.r.t Online Security



```
color_map = {"Yes": '#FFA15A', "No": '#00CC96'}  
fig = px.histogram(df, x="Churn", color="PaperlessBilling", title="Churn distribution w.r.t. Paperless Billing", color_discrete  
fig.update_layout(width=700, height=500, bargap=0.1)  
fig.show()
```




Chrun distribution w.r.t. Paperless Billing



```
fig = px.histogram(df, x="Churn", color="TechSupport", barmode="group", title="Chrun distribution w.r.t. TechSupport")  
fig.update_layout(width=700, height=500, bargap=0.1)  
fig.show()
```




Chrun distribution w.r.t. TechSupport




```
color_map = {"Yes": '#00CC96', "No": '#B6E880'}
fig = px.histogram(df, x="Churn", color="PhoneService", title="Chrun distribution w.r.t. Phone Service", color_discrete_map=colo
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```




Chrun distribution w.r.t. Phone Service



```
sns.set_context("paper", font_scale=1.1)
ax = sns.kdeplot(df.MonthlyCharges[(df["Churn"] == 'No') ],
                 color="Red", shade = True);
ax = sns.kdeplot(df.MonthlyCharges[(df["Churn"] == 'Yes') ],
                 ax =ax, color="Blue", shade= True);
ax.legend(["Not Churn", "Churn"], loc='upper right');
ax.set_ylabel('Density');
ax.set_xlabel('Monthly Charges');
ax.set_title('Distribution of monthly charges by churn');
```



```
ax = sns.kdeplot(df.TotalCharges[(df["Churn"] == 'No') ],
                 color="Gold", shade = True);
ax = sns.kdeplot(df.TotalCharges[(df["Churn"] == 'Yes') ],
                 ax=ax, color="Green", shade= True);
ax.legend(["Not Churn", "Churn"], loc='upper right');
ax.set_ylabel('Density');
ax.set_xlabel('Total Charges');
ax.set_title('Distribution of total charges by churn');
```



```
fig = px.box(df, x='Churn', y = 'tenure')


# Update yaxis properties
fig.update_yaxes(title_text='Tenure (Months)', row=1, col=1)
# Update xaxis properties
fig.update_xaxes(title_text='Churn', row=1, col=1)

# Update size and title
fig.update_layout(autosize=True, width=750, height=600,
                  title_font=dict(size=25, family='Courier'),
                  title='<b>Tenure vs Churn</b>',
)

fig.show()
```



Tenure vs Churn




```
plt.figure(figsize=(25, 10))
```

```
corr = df.apply(lambda x: pd.factorize(x)[0]).corr()
```

```
mask = np.triu(np.ones_like(corr, dtype=bool))
```

```
ax = sns.heatmap(corr, mask=mask, xticklabels=corr.columns, yticklabels=corr.columns, annot=True, linewidths=.2, cmap='coolwarm', vmin=
```



```
def object_to_int(dataframe_series):
    if dataframe_series.dtype=='object':
        dataframe_series = LabelEncoder().fit_transform(dataframe_series)
    return dataframe_series
```

```
df = df.apply(lambda x: object_to_int(x))
df.head()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	D
0	0	0	1	0	1	0	1	0	0	2	
1	1	0	0	0	34	1	0	0	2	0	
2	1	0	0	0	2	1	0	0	2	2	
3	1	0	0	0	45	0	1	0	2	0	
4	0	0	0	0	2	1	0	1	0	0	

Next steps:

[Generate code with df](#)[View recommended plots](#)

```
plt.figure(figsize=(14,7))
df.corr()['Churn'].sort_values(ascending = False)
```

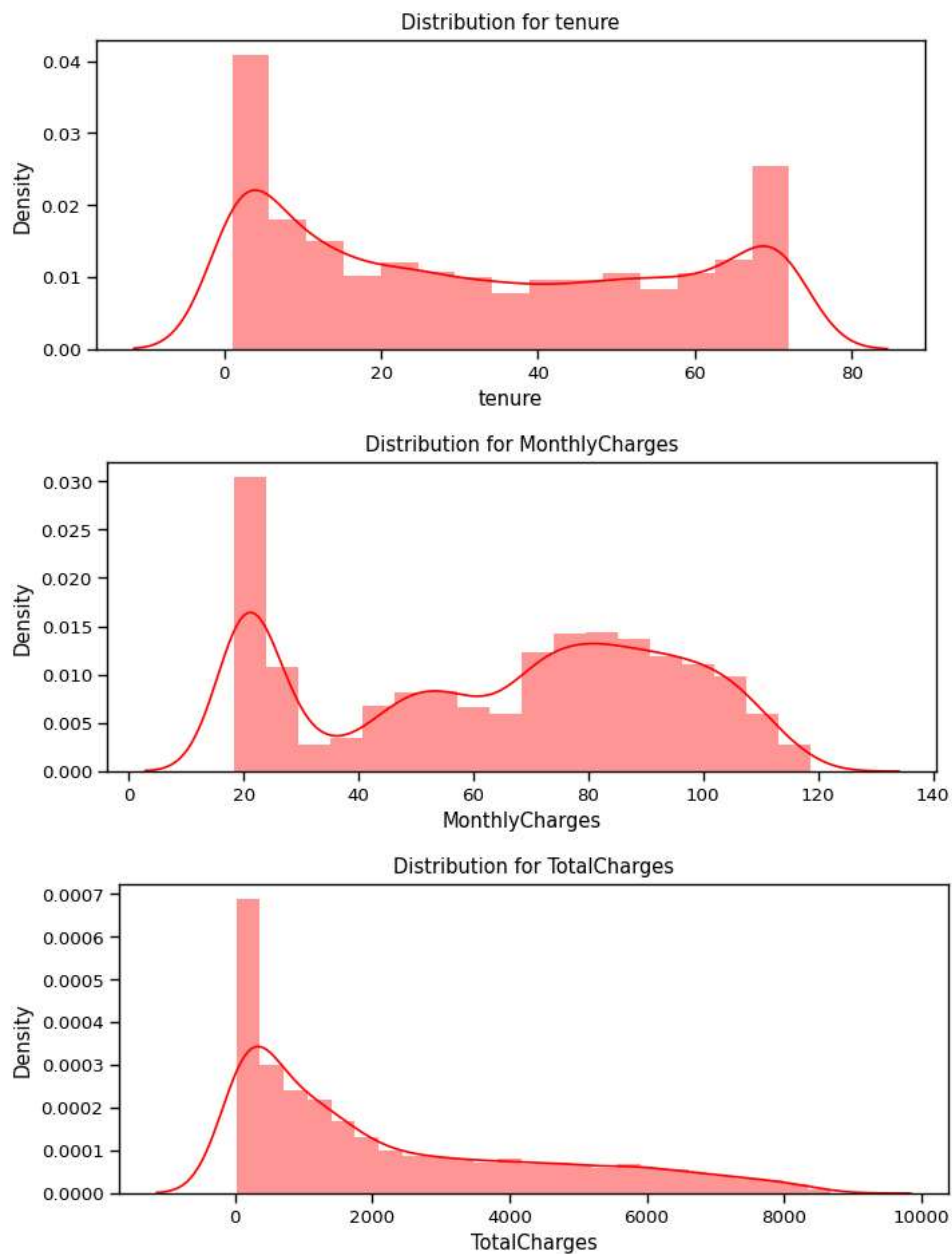
```
Churn      1.000000
MonthlyCharges  0.192858
PaperlessBilling  0.191454
SeniorCitizen  0.150541
PaymentMethod  0.107852
MultipleLines  0.038043
PhoneService   0.011691
gender         -0.008545
StreamingTV    -0.036303
StreamingMovies -0.038802
InternetService -0.047097
Partner        -0.149982
Dependents     -0.163128
DeviceProtection -0.177883
OnlineBackup   -0.195290
TotalCharges   -0.199484
TechSupport    -0.282232
OnlineSecurity -0.289050
tenure         -0.354049
Contract       -0.396150
Name: Churn, dtype: float64
<Figure size 1400x700 with 0 Axes>
```

```
X = df.drop(columns = ['Churn'])
y = df['Churn'].values
```




```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.30, random_state = 40, stratify=y)
```

```
def distplot(feature, frame, color='r'):
    plt.figure(figsize=(8,3))
    plt.title("Distribution for {}".format(feature))
    ax = sns.distplot(frame[feature], color= color)
```

```
num_cols = ["tenure", 'MonthlyCharges', 'TotalCharges']
for feat in num_cols: distplot(feat, df)
```



```
df_std = pd.DataFrame(StandardScaler().fit_transform(df[num_cols].astype('float64')),  
                      columns=num_cols)  
for feat in numerical_cols: distplot(feat, df_std, color='c')
```

```
# Divide the columns into 3 categories, one for standardisation, one for label encoding and one for one hot encoding
```

```
cat_cols_ohe = ['PaymentMethod', 'Contract', 'InternetService'] # those that need one-hot encoding
cat_cols_le = list(set(X_train.columns) - set(num_cols) - set(cat_cols_ohe)) #those that need label encoding
```

```
scaler= StandardScaler()
```

```
X_train[num_cols] = scaler.fit_transform(X_train[num_cols])
X_test[num_cols] = scaler.transform(X_test[num_cols])
```

```
knn_model = KNeighborsClassifier(n_neighbors = 11)
knn_model.fit(X_train,y_train)
predicted_y = knn_model.predict(X_test)
accuracy_knn = knn_model.score(X_test,y_test)
print("KNN accuracy:",accuracy_knn)
```



```
KNN accuracy: 0.7758293838862559
```

```
print(classification_report(y_test, predicted_y))
```



	precision	recall	f1-score	support
0	0.83	0.87	0.85	1549
1	0.59	0.52	0.55	561
accuracy			0.78	2110
macro avg	0.71	0.69	0.70	2110
weighted avg	0.77	0.78	0.77	2110

```

svc_model = SVC(random_state = 1)
svc_model.fit(X_train,y_train)
predict_y = svc_model.predict(X_test)
accuracy_svc = svc_model.score(X_test,y_test)
print("SVM accuracy is :",accuracy_svc)

```

SVM accuracy is : 0.8075829383886256

```
print(classification_report(y_test, predict_y))
```

```

precision    recall  f1-score   support

      0       0.84      0.92      0.88       1549
      1       0.69      0.50      0.58        561

 accuracy          0.81       2110
  macro avg       0.76      0.71      0.73       2110
 weighted avg     0.80      0.81      0.80       2110

```

```

model_rf = RandomForestClassifier(n_estimators=500 , oob_score = True, n_jobs = -1,
                                random_state =50, max_features = "auto",
                                max_leaf_nodes = 30)

model_rf.fit(X_train, y_train)

```

```

# Make predictions
prediction_test = model_rf.predict(X_test)
print (metrics.accuracy_score(y_test, prediction_test))

```

0.8137440758293839

```
print(classification_report(y_test, prediction_test))
```

```

precision    recall  f1-score   support

      0       0.84      0.92      0.88       1549
      1       0.71      0.51      0.59        561

 accuracy          0.81       2110
  macro avg       0.77      0.72      0.74       2110
 weighted avg     0.80      0.81      0.80       2110

```

```


plt.figure(figsize=(4,3))
sns.heatmap(confusion_matrix(y_test, prediction_test),
            annot=True,fmt = "d",linecolor="k",linewidths=3)

```

```

plt.title(" RANDOM FOREST CONFUSION MATRIX",fontsize=14)
plt.show()


```



```

y_rfpred_prob = model_rf.predict_proba(X_test)[:,-1]
fpr_rf, tpr_rf, thresholds = roc_curve(y_test, y_rfpred_prob)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_rf, tpr_rf, label='Random Forest',color = "r")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Random Forest ROC Curve',fontsize=16)
plt.show();

```



```
lr_model = LogisticRegression()
lr_model.fit(X_train,y_train)
accuracy_lr = lr_model.score(X_test,y_test)
print("Logistic Regression accuracy is :",accuracy_lr)
```

Logistic Regression accuracy is : 0.8090047393364929

```
lr_pred= lr_model.predict(X_test)
report = classification_report(y_test,lr_pred)
print(report)
```




	precision	recall	f1-score	support
0	0.86	0.89	0.87	1549
1	0.66	0.58	0.62	561
accuracy			0.81	2110
macro avg	0.76	0.74	0.75	2110
weighted avg	0.80	0.81	0.81	2110

```
plt.figure(figsize=(4,3))
sns.heatmap(confusion_matrix(y_test, lr_pred),
            annot=True,fmt = "d",linecolor="k",linewidths=3)
```

```
plt.title("LOGISTIC REGRESSION CONFUSION MATRIX",fontsize=14)
plt.show()
```




LOGISTIC REGRESSION CONFUSION MATRIX



```
y_pred_prob = lr_model.predict_proba(X_test)[:,:1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr, tpr, label='Logistic Regression',color = "r")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Logistic Regression ROC Curve',fontsize=16)
plt.show();
```



Logistic Regression ROC Curve



```
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train,y_train)
predictdt_y = dt_model.predict(X_test)
accuracy_dt = dt_model.score(X_test,y_test)
print("Decision Tree accuracy is :",accuracy_dt)
```



Decision Tree accuracy is : 0.7303317535545024

```
print(classification_report(y_test, predictdt_y))
```



	precision	recall	f1-score	support
0	0.82	0.81	0.81	1549
1	0.49	0.52	0.51	561
accuracy			0.73	2110
macro avg	0.66	0.66	0.66	2110
weighted avg	0.74	0.73	0.73	2110